

VI Accumulators

When you ask ISL+ to apply some function *f* to an argument *a*, you usually get some value *v*. If you evaluate `(f a)` again, you get *v* again. As a matter of fact, you get the same result no matter how often you request the evaluation of `(f a)`. Whether the function is applied for the first time or the hundredth time, whether the application is located in DrRacket’s interactions area or inside the function itself, doesn’t matter. The function works according to its purpose statement, and that’s all you need to know.

This principle of context-independence plays a critical role in the design of recursive functions. When it comes to coding, you are free to assume that the function computes what the purpose statement promises—even if the function isn’t defined yet. In particular, you are free to use the results of recursive calls to create the code of some function, usually one of its `cond` clauses. The template and coding steps of the design recipes for both structurally and generative recursive functions rely on this idea.

Although context-independence facilitates the design of functions, it also causes two problems. The general idea is that context-independence induces a loss of knowledge during a recursive evaluation; a function does not “know” whether it is called on a complete list or on a piece of that list. For structurally recursive programs this loss of knowledge means that they may have to traverse data more than once, inducing a grave performance cost. For functions that employ generative recursion, the loss means that the function may not be able to compute the result; instead the function loops forever for certain inputs. The preceding part illustrates this second problem with a graph traversal function that cannot find a path between two nodes for a circular graph.

This part introduces a variant of the design recipes to address this “loss of context” problem. Since we wish to retain the principle that `(f a)` returns the same result no matter how often it is evaluated, our only solution is to add **an argument that represents the context** of the function call. We call this additional argument an *accumulator*. During the traversal of data, the recursive calls continue to receive new regular arguments while accumulators change in relation to the other arguments and the context of the call.

Designing functions with accumulators correctly is clearly more complex than any of the design approaches from the preceding chapters. The key is to understand the relationship between the proper arguments and the accumulators. The following chapters explain how to design functions with accumulators and how they work.

The function application may also loop forever or signal an error, but let’s ignore these possibilities for now. We also ignore random, which is the only real exception to this rule.

36 The Loss of Knowledge


Both functions designed according to structural recipes and the generative one suffer from the loss of knowledge, though in different ways. This chapter explains with two examples—one from each category—how the lack of contextual knowledge affects the performance of functions. While the first section is about structural recursion, the second one addresses concerns in the generative realm.

36.1 A Problem with Structural Processing


Let's start with a seemingly straightforward example:

Sample Problem: You are working for a geometer team that will measure the length of roads segments. The team asked you to design a program that translates these relative distances between a series of road points into absolute distances for some starting point.

For example, we might be given a line such as this:



Each number specifies the distance between two dots. What we need is the following picture, where each dot is annotated with the distance to the left-most end:



Designing a program that performs this calculation is at this point an exercise in structural function design. Figure 103 contains the complete program. When the given list is not '(), the natural recursion computes the absolute distance of the remainder of the dots to the first item on (rest alon). Because the first item is not the actual origin and has a distance of (first alon) to the origin, we must add (first alon) to each number on the result of the natural recursion. This second step—adding a number to each item on a list of numbers—requires an auxiliary function.

```
; [List-of Number] -> [List-of Number]
; convert a list of relative distances to a list of absolute distances
; the first item on the list represents the distance to the origin

(check-expect (relative->absolute '(50 40 70 30 30))
              '(50 90 160 190 220))

(define (relative->absolute l)
  (cond
    [(empty? l) '()]
    [else (local ((define rest-of-l (relative->absolute (rest l)))
                  (define adjusted (add-to-each (first l) rest-of-l)))
            (cons (first l) adjusted))]))

; Number [List-of Number] -> [List-of Number]
; add n to each number on alon

(check-expect (cons 50 (add-to-each 50 '(40 110 140 170)))
              '(50 90 160 190 220))

(define (add-to-each n alon)
  (cond
    [(empty? alon) '()]
    [else (cons (+ (first alon) n) (add-to-each n (rest alon)))]))
```

Figure 103: Converting relative distances to absolute distances

While designing the program is relatively straightforward, using it on larger and larger lists reveals a problem. Consider the evaluation of the following expression:

```
(relative->absolute (list 0 ... size))
```

Use

As we increase `size`, the time needed grows even faster:

size	1000	2000	3000	4000	5000	6000	7000
time	25	109	234	429	689	978	1365

Instead of doubling as we go from 1000 to 2000 items, the time quadruples. This is also the approximate relationship for going from 2000 to 4000, and so on.

Exercise 401. Reformulate `add-to-each` using `map` and `lambda`. ■

Exercise 402. Determine the abstract running time of `relative->absolute`.

Hint Evaluate the expression

```
(relative->absolute (list 0 ... size))
```

by hand. Start by replacing `size` with 1, 2, and 3. How many natural recursions of `relative->absolute` and `add-to-each` are required each time? ■

Considering the simplicity of the problem, the amount of “work” that the program performs is surprising. If we were to convert the same list by hand, we would tally up the total distance and just add it to the relative distances as we take another step along the line. Why can’t a program use this idea?

Let’s attempt to design a second version of the function that is closer to our manual method. The new function is still a list-processing function, so we start from the appropriate template:

```
(define (relative->absolute/a alon)
  (cond
    [(empty? alon) ...]
    [else ... (first alon) ... (relative->absolute/a (rest alon)) ...]))
```

Now imagine an “evaluation” of `(relative->absolute/a (list 3 2 7))`:

```
= (cons ... 3 ...
   (convert (list 2 7)))

= (cons ... 3 ...
   (cons ... 2 ...
    (convert (list 7))))

= (cons ... 3 ...
   (cons ... 2 ...
    (cons ... 7 ...
     (convert '()))))
```

The first item of the result list should obviously be 3, and it is easy to construct this list. But, the second one should be `(+ 3 2)`, yet the second instance of `relative->absolute/a` has no way of “knowing” that the first item of the **original** list is 3. The “knowledge” is lost.

Put differently, the problem is that recursive functions are independent of their context. A function processes `L` in `(cons N L)` in the same manner as in `(cons K L)`. Indeed, it would also process `L` in that manner if it were given `L` by itself.

To make up for the loss of “knowledge,” we equip the function with an additional parameter: `accu-dist`. The new parameter represents the accumulated distance, which is the tally that we keep when we convert a list of relative distances to a list of absolute

(build-
list
size
add1)
to
construct
these
The
time
of
evaluation
will
differ
from
computer
to
computer
and
from
year
to
year.
These
measurements
were
conducted
in
2014
on a
MacMini
running
OS X
10.8.5;
the
previous
measurement
took
place
in
1998,
and
the
times
were
100x
larger.

distances. Its initial value must be 0. As the function processes the numbers on the list, it must add them to the tally.

Here is the revised definition:

```
(define (relative->absolute/a alon accu-dist)
  (cond
    [(empty? alon) '()]
    [else (local ((define tally (+ (first alon) accu-dist)))
              (cons tally (relative->absolute/a (rest alon) tally))))])
```

The recursive application consumes the rest of the list and the new absolute distance of the current point to the origin. Although this means that two arguments are changing simultaneously, the change in the second one strictly depends on the first argument. The function is still a plain list-processing procedure.

Evaluating our running example `(relative->absolute/a (list 3 2 7))` again, shows how much the use of an accumulator simplifies the conversion process:

```
= (relative->absolute/a (list 3 2 7) 0)
= (cons 3 (relative->absolute/a (list 2 7) 3))
= (cons 3 (cons 5 (relative->absolute/a (list 7) 5)))
= (cons 3 (cons 5 (cons 12 (relative->absolute/a '() 12))))
= (cons 3 (cons 5 (cons 12 '()))))
```

Each item in the list is processed once. When `relative->absolute/a` reaches the end of the argument list, the result is completely determined and no further work is needed. In general, the function performs on the order of N natural recursion steps for a list with N items.

One minor problem with the new definition is that unlike `relative->absolute`, the new function consumes two arguments not just one. Worse, someone might accidentally misuse `relative->absolute/a` by applying it to a list of numbers and a number that isn't 0. We can solve both problems with a function definition that uses a `local` definition to encapsulate `relative->absolute/a`; figure 104 shows the result. Now, `relative->absolute` and `relative->absolute.v2` are indistinguishable with respect to the input-output relationship.

```
; [List-of Number] -> [List-of Number]
; convert a list of relative distances to a list of absolute distances
; the first item on the list represents the distance to the origin

(check-expect (relative->absolute.v2 '(50 40 70 30 30))
              '(50 90 160 190 220))

(define (relative->absolute.v2 alon0)
  (local (; [List-of Number] Number -> [List-of Number]
          (define (relative->absolute/a alon accu-dist)
            (cond
              [(empty? alon) '()]
              [else
               (local ((define accu (+ (first alon) accu-dist)))
                 (cons accu (relative->absolute/a (rest alon) accu))))]))
    (relative->absolute/a alon0 0)))
```

Figure 104: Converting relative distances with an accumulator

Now let's look at how this version of the program performs. To this end, we evaluate

```
(relative->absolute.v2 (list 0 ... size))
```

and tabulate the results for several values of `size`:

size	1000	2000	3000	4000	5000	6000	7000
time	0	0	0	0	0	1	1

Amazingly `relative->absolute.v2` never takes more than one second to process such lists, even for a list of 7000 numbers. Comparing this performance to the one of `relative->absolute`, you may think that accumulators are a miracle cure for all slow-running programs. Unfortunately, this isn't the case, but when a structurally recursive functions has to re-process the result of the natural recursion you should definitely consider the use of accumulators.

Exercise 403. With a bit of design and a bit of tinkering a friend of yours came up with the following solution for the sample problem:

```
; [List-of Number] -> [List-of Number]
; convert a list of relative distances to a list of absolute distances
; the first item on the list represents the distance to the origin

(check-expect (relative->absolute '(50 40 70 30 30))
               '(50 90 160 190 220))

(define (relative->absolute l)
  (reverse
   (foldr (lambda (f l) (cons (+ f (first l)) l))
         (list (first l))
         (reverse (rest l))))))
```

Messrs.
Adrian
German
and
Mardin
Yadegar
suggested
this
exercise.

This simple solution merely uses well-known ISL+ functions: `reverse` and `foldr`. Using `lambda`, as you know, is just a convenience. You may also recall from [Abstraction](#) that `foldr` is designable with the design recipes presented in the first two parts of the book.

Does your friend's solution mean there is no need for our complicated design in this motivational section? xFor a solution, see [Recognizing the Need for an Accumulator](#), but do reflect on the question first and better still try to design `reverse` on your own. ■


36.2 A Problem with Generative Recursion

Let us revisit the problem of “traveling” along a path in a graph:

Sample Problem: Design an algorithm that checks whether two nodes are connected in a simple graph. In a *simple graph*, each node has exactly one, one-directional connection to another node, possibly itself.

[Algorithms that Backtrack](#) covered the variant where the algorithm has to discover the path. This sample problem is simpler than that, because this section focuses on the design of an accumulator version of the algorithm.

Consider the sample graph in [figure 105](#). There are six nodes: *A* through *F*, and six connections. To get from *A* to *E*, you must go through *B* and *C*. It is impossible, though, to reach *F* from *A* or from any other node besides *F* itself.



```

(define a-simple-graph
  '((A B)
    (B C)
    (C E)
    (D E)
    (E B)
    (F F)))

```

Figure 105: A simple graph

The right part of figure 105 shows how to represent this graph with nested lists. Each node is represented by a list of two symbols. The first symbol is the label of the node; the second one is the single node that is reachable from the first one. Here are the relevant data definitions:

```

; A SimpleGraph is a [List-of Connection]
; A Connection is (list Node Node)
; A Node is a Symbol

```

They are straightforward translations of our informal descriptions.

We already know that the problem calls for generative recursion, and it is easy to create the header material:

```

; Node Node SimpleGraph -> Boolean
; is there a path from origination to destination in sg

(check-expect (path-exists? 'A 'E a-simple-graph) #true)
(check-expect (path-exists? 'A 'F a-simple-graph) #false)

(define (path-exists? origination destination sg)
  #false)

```

What we need are answers to the four basic questions of the recipe for generative recursion:

- The problem is trivial if the nodes `origination` and `destination` are the same.
- The trivial solution is `#true`.
- If `origination` is not the same as `destination`, there is only one thing we can do: step to the immediate neighbor and search for `destination` from there.
- There is no need to do anything if we find the solution to the new problem. If `origination`'s neighbor is connected to `destination`, then so is `origination`. Otherwise there is no connection.

From here we just need to express these answers in ISL+ to obtain a full-fledged program.

```

; Node Node SimpleGraph -> Boolean
; is there a path from origination to destination in sg

(check-expect (path-exists? 'A 'E a-simple-graph) #true)
(check-expect (path-exists? 'A 'F a-simple-graph) #false)

```

```

(define (path-exists? origination destination sg)
  (cond
    [(symbol=? origination destination) #t]
    [else (path-exists? (neighbor origination sg) destination sg)]))

; Node SimpleGraph -> Node
; determine the node that is connected to a-node in sg

(check-expect (neighbor 'A a-simple-graph) 'B)
(check-error (neighbor 'G a-simple-graph) "neighbor: not a node")

(define (neighbor a-node sg)
  (cond
    [(empty? sg) (error "neighbor: not a node")]

    [else (if (symbol=? (first (first sg)) a-node)
              (second (first sg))
              (neighbor a-node (rest sg)))]))

```

Figure 106: Finding a path in a simple graph

Figure 106 contains the complete program, including the function for looking up the neighbor of a node in a simple graph—a straightforward exercise in structural recursion—and test cases for both. Don’t run the program, however. If you do, be ready with your mouse to stop the run-away program. Indeed, even a casual look at the function suggests that we have a problem. Although the function is supposed to produce `#false` if there is no path from `origination` to `destination`, the program doesn’t contain `#false` anywhere. Conversely, we need to ask what the function actually does when there is no path between two nodes.

Take another look at figure 105. In this simple graph there is no path from *C* to *D*. The connection that leaves *C* passes right by *D* and instead goes to *E*. So let’s look at how

```
(path-exists? 'C 'D '((A B) (B C) (C E) (D E) (E B) (F F)))
```

is evaluated:

```

= (path-exists? 'E 'D '((A B) (B C) (C E) (D E) (E B) (F F)))
= (path-exists? 'B 'D '((A B) (B C) (C E) (D E) (E B) (F F)))
= (path-exists? 'C 'D '((A B) (B C) (C E) (D E) (E B) (F F)))

```

The hand-evaluation confirms that as the function recurs, it calls itself with the exact same arguments again and again. In other words, the evaluation never stops.

Our problem with `path-exists?` is again a loss of “knowledge,” similar to that of `relative->absolute` in the preceding section. Like `relative->absolute`, the design of `path-exists?` uses a recipe and assumes context-independence for recursive calls. In the case of `path-exists?` this means, in particular, that the function doesn’t “know” whether a previous application in the current chain of recursions received the exact same arguments.

The solution to this design problem follows the pattern of the preceding section. We add a parameter, which we call `seen` and which represents the accumulated list of origination nodes that the function has encountered, starting with the original application. Its initial value must be `'()`. As the function checks on a specific origination and moves to its neighbors, origination is added to `seen`.

Here is a first revision of `path-exists?`, dubbed `path-exists?/a`:

```

; Node Node SimpleGraph [List-of Node] -> Boolean
; is there a path from origination to destination in sg
; assume the nodes in seen are known not to solve the problem
(define (path-exists?/a origination destination sg seen)
  (cond
    [(symbol=? origination destination) #t]
    [else (path-exists?/a (neighbor origination sg) destination sg
                          (cons origination seen))]))

```

The addition of the new parameter alone does not solve our problem, but, as the hand-evaluation of

```
(path-exists?/a 'C 'D '((A B) (B C) (C E) (D E) (E B) (F F)) '())
```

shows, provides the foundation for one:

```

= (path-exists?/a 'E 'D '((A B) (B C) (C E) (D E) (E B) (F F)) '(C))
= (path-exists?/a 'B 'D '((A B) (B C) (C E) (D E) (E B) (F F)) '(E C))
= (path-exists?/a 'C 'D '((A B) (B C) (C E) (D E) (E B) (F F)) '(B E C))

```

In contrast to the original function, the revised function no longer calls itself with the exact same arguments. While the three arguments proper are again the same for the third recursive application, the accumulator argument is different from that of the first application. Instead of '(), it is now '(B E C). The new value represents the fact that during the search of a path from 'C to 'D, the function has inspected 'B, 'E, and 'C as starting points.

All we need to do now, is to make the algorithm exploit the accumulated knowledge. Specifically, the algorithm can determine whether the given origination is already an item in seen. If so, the problem is also trivially solvable yielding #false as the solution. Figure 107 contains the definition of path-exists.v2?, which is the revision of path-exists?. The definition refers to member?, an ISL+ function.

```

; Node Node SimpleGraph -> Boolean
; is there a path from origination to destination in sg

(check-expect (path-exists.v2? 'A 'E a-simple-graph) #true)
(check-expect (path-exists.v2? 'A 'F a-simple-graph) #false)

(define (path-exists.v2? origination destination sg)
  (local (; Node Node SimpleGraph [List-of Node] -> Boolean
        (define (path-exists?/a origination seen)
          (cond
            [(symbol=? origination destination) #t]
            [(member? origination seen) #f]
            [else (path-exists?/a (neighbor origination sg)
                                  (cons origination seen))]))
    (path-exists?/a origination '())))

```

Figure 107: Finding a path in a simple graph with an accumulator

The definition of path-exists.v2? also eliminates the two minor problems with the first revision. By localizing the definition of the accumulating function, we can ensure that the first call always uses '() as the initial value for seen. And, path-exists.v2? satisfies the exact same contract and purpose statement as the path-exists? function.

Still, there is a significant difference between path-exists.v2? and relative-to-

`absolute2`. Whereas the latter was equivalent to the original function, `path-exists.v2?` improves on `path-exists?`. While the latter fails to find an answer for some inputs, `path-exists.v2?` finds a solution for any simple graph.

Exercise 404. Modify the definitions of `find-path` and `find-path/list` in [figure 97](#) so that they produce `#false`, even if they encounter the same starting point twice. ■

37 Designing Accumulator-Style Functions

The preceding chapter illustrates the need for accumulating extra knowledge with two examples. In one case, accumulation makes it easy to understand the function and yields one that is far faster than the original version. In the other case, accumulation is necessary for the function to work properly. In both cases though, the need for accumulation becomes only apparent once a properly designed function exists.

Generalizing from the preceding chapter suggests that the design of accumulator functions has two major aspects:

1. the recognition that a function benefits from an accumulator;
2. an understanding of what the accumulator represents with respect to the design.

The first two sections address these two questions. Because the second one is a difficult topic, the third section illustrates it with a series of examples that convert regular functions into accumulating ones.

37.1 Recognizing the Need for an Accumulator

Recognizing the need for accumulators is not an easy task. We have seen two reasons, and they are the most prevalent reasons for adding accumulator parameters. In either case, it is critical that we first built a complete function based on a design recipe. Then we study the function and look for one of the following characteristics:

1. If a structurally recursive function processes the result of its natural recursion with an auxiliary, recursive function, consider the use of an accumulator parameter.

Take a look at the definition of `invert`:

```
; [List-of X] -> [List-of X]
; construct the reverse of alox

(check-expect (invert '(a b c)) '(c b a))

(define (invert alox)
  (cond
    [(empty? alox) '()]
    [else (add-as-last (first alox) (invert (rest alox)))]))

; X [List-of X] -> [List-of X]
; add an-x to the end of alox

(check-expect (add-as-last 'a '(c b)) '(c b a))

(define (add-as-last an-x alox)
  (cond
```

```
[(empty? alox) (list an-x)]
[else (cons (first alox) (add-as-last an-x (rest alox))))])
```

The result of the recursive application produces the reverse of the rest of the list. It is processed by `add-as-last`, which adds the first item to the reverse of the rest and thus creates the reverse of the entire list. This second, auxiliary function is also recursive. We have thus identified a potential candidate.

It is now time to study some hand-evaluations, as we did in [A Problem with Structural Processing](#), to see whether an accumulator helps. Consider the following expression:

```
(invert '(a b c))
```

Here is how you calculate how `invert` determines the result when given `'(a b c)`:

```
= (add-as-last 'a (invert '(b c)))
= (add-as-last 'a (add-as-last 'b (invert '(c))))
= (add-as-last 'a (add-as-last 'b (add-as-last 'c (invert '()))))
= (add-as-last 'a (add-as-last 'b (add-as-last 'c '())))
= (add-as-last 'a (add-as-last 'b '(c)))
= (add-as-last 'a '(c b))
= '(c b a)
```

Eventually `invert` reaches the end of the given list—just like `add-as-last`—and if it knew which items to put there, there would be no need for the auxiliary function.

2. If we are dealing with a function based on generative recursion, we are faced with a much more difficult task. Our goal must be to understand whether the algorithm can fail to produce a result for inputs for which we expect a result. If so, adding a parameter that accumulates knowledge may help. Because these situations are complex, we defer the discussion of an example to [More Uses of Accumulation](#).

Exercise 405. Does the insertion `sort>` function from [Recursive Auxiliary Functions](#) need an accumulator? If so, why? If not, why not? ■

37.2 Adding Accumulators

Once you have decided that an existing function should be equipped with an accumulator, take these two steps:

- Determine the knowledge that the accumulator represents, what kind of data to use, and how the knowledge is acquired as data.

For example, for the conversion of relative distances to absolute distances, it suffices to accumulate the total distance encountered so far. As the function processes the list of relative distances, it adds each new relative distance found to the accumulator's current value. For the routing problem, the accumulator remembers every node encountered. As the path-checking function traverses the graph, it `conses` each new node on to the accumulator.

In general, you want to proceed as follows.

1. Create an accumulator template:

```
; Domain -> Range
(define (function d0)
  (local (; Domain AccumulatorDomain -> Range
```

```

; accumulator ...
(define (function/a d a)
  ...)
(function/a d0 a0))

```

Sketch a manual evaluation of an application of `function` to understand the nature of the accumulator.

- Determine the kind of data that the accumulator tracks.

Write down a statement that explains the accumulator as a relationship between the argument `d` of the auxiliary `function/a` and the original argument `d0`.

Note The relationship remains constant—also called **invariant**—over the course of the evaluation. Because of this property, an accumulator statement is also called an accumulator *invariant*.

- Use the accumulator statement to determine the initial value `a0` for `a`.
 - Also exploit the accumulator statement to determine how to compute the accumulator for the recursive function calls within the definition of `function/a`.
- Exploit the accumulator's knowledge for the design of the auxiliary function.

For a structurally recursive function, the accumulator's value is typically used in the base case, that is, the `cond` clause that does not recur. For functions that use generative recursive functions, the accumulated knowledge might be used in an existing base case, in a new base case, or in the `cond` clauses that deal with generative recursion.

As you can see, the key is the precise description of the role of the accumulator. It is therefore important to practice this skill.

Let's take a look at the `invert` example:

```

; [List-of X] -> [List-of X]
; construct the reverse of alox0

(check-expect (invert.v2 '(a b c)) '(c b a))

(define (invert.v2 alox0)
  (local (; [List-of X] ??? -> [List-of X]
          ; construct the reverse of alox
          ; accumulator ...

          (define (invert/a alox a)
            (cond
              [(empty? alox) ...]
              [else (invert/a (rest alox) ... a ...)])))
    (invert/a alox0 ...)))

```

As illustrated in the preceding section, this template suffices to sketch a manual evaluation of an expression such as

```
(invert '(a b c))
```

Here is the idea:

```

= (invert/a '(a b c) a0)
= (invert/a '(b c) ... 'a ... a0)

```

```
= (invert/a '(c) ... 'b ... 'a ... a0)
= (invert/a '() ... 'c ... 'b ... 'a ... a0)
```

This sketch suggests that `invert/a` can keep track of all the items it has seen in a list that tracks the difference between `alox0` and `a` in reverse order. The initial value is clearly `'()`; updating the accumulator inside of `invert/a` with `cons` produces exactly the desired value when `invert/a` reaches `'()`.

Here is a refined template that includes these insights:

```
(define (invert.v2 alox0)
  (local (; [List-of X] [List-of X] -> [List-of X]
          ; construct the reverse of alox
          ; accumulator a is the list of all those items
          ; on alox0 that precede alox in reverse order
          (define (invert/a alox a)
            (cond
              [(empty? alox) a]
              [else (invert/a (rest alox) (cons (first alox) a))]))
    (invert/a alox0 '())))
```

While the body of the `local` definition initializes the accumulator with `'()`, the recursive call uses `cons` to add the current head of `alox` to the accumulator. In the base case, `invert/a` uses the knowledge in the accumulator—the reversed list.

Note how once again `invert.v2` traverses the list just. In contrast, `invert` re-processes every result of its natural recursion with `add-as-last`. Stop! Measure how much faster `invert.v2` runs than `invert` on the same list.

Terminology Programmers use the phrase *accumulator-style function* when they discuss functions that use an accumulator parameter. Examples of functions in accumulator-style are `relative->absolute/a`, `path-exists?/a`, and `invert/a`.

37.3 Transforming Functions into Accumulator-Style

Articulating the accumulator statement is difficult but without formulating a good invariant, it is impossible to understand an accumulator-style function. Since the goal of a programmer is to make sure that others who follow understand the code easily, practicing this skill is critical. And formulating invariants is deserves a lot of practice.

The goal of this section is to study the formulation of accumulator statements with three case studies: a summation function, the factorial function, and a tree-traversal function. Each such case is about the conversion of a structurally recursive function into accumulator style. None actually call for the use of an accumulator parameter. But they are easily understood and, with the elimination of all other distractions, using such examples allows us to focus on the articulation of the accumulator invariant.

For the first example, consider the following definition of the `sum` function:

```
; [List-of Number] -> Number
; compute the sum of the numbers on alon

(check-expect (sum '(10 4 6)) 20)

(define (sum alon)
  (cond
```

```

[(empty? alon) 0]
[else (+ (first alon) (sum (rest alon)))])

```

Here is the first step toward an accumulator version:

```

; [List-of Number] -> Number
; compute the sum of the numbers on alon0

(check-expect (sum.v2 '(10 4 6)) 20)

(define (sum.v2 alon0)
  (local (; [List-of Number] ??? -> Number
          ; compute the sum of the numbers on alon
          ; accumulator ...
          (define (sum/a alon a)
            (cond
              [(empty? alon) ...]
              [else (... (sum/a (rest alon) ... a ...) ...) ])))
    (sum/a alon0 ...)))

```

As suggested by our first step, we have put the template for `sum/a` into a `local` definition, added an accumulator parameter, and renamed `sum` `'s` parameter.

Here are two side-by-side sketches of hand evaluations:

<pre> (sum '(10 4)) = = (+ 10 (sum '(4))) = (+ 10 (+ 4 (sum '()))) = (+ 10 (+ 4 (+ 0))) ... = 20.0 </pre>	<pre> (sum.v2 '(10 4)) = = (sum/a '(10 4) a0) = (sum/a '(4) ... 10 ... a0) = (sum/a '() ... 4 ... 10 ... a0) ... = 20.0 </pre>
---	--

A comparison immediately suggests the central idea, namely, that `sum/a` can use the accumulator to add up the numbers as it encounters.

Concerning the accumulator invariant, this analysis suggest `a` represents the sum of the numbers encountered so far:

`a` represents the sum of the numbers that `alon` lacks in comparison to `alon0`

For example, if

```

alon0 = '(10 4 6)
alon  = '(6)

```

the invariant forces the accumulator's value to be `14`. In contrast, when

```

alon0 = '(10 4 6)
alon  = '()

```

`a` must be `20`.

Given this precise invariant, the rest of the design is straightforward again:

```

(define (sum.v2 alon0)
  (local (; [List-of Number] ??? -> Number
          ; compute the sum of the numbers on alon
          ; accumulator a represents the sum of the numbers
          ; that alon lacks in comparison to alon0

```

```

(define (sum/a alon a)
  (cond
    [(empty? alon) a]
    [else (sum/a (rest alon) (+ (first alon) a))]))

(sum/a alon0 0))

```

If `alon` is `'()`, `sum/a` returns `a` because it represents the sum of all numbers on `alon`. The invariant also implies that `0` is the initial value for `a0` and `+` updates the accumulator by adding the number that is about to be “forgotten”—`(first alon)`—to the accumulator `a`.

Exercise 406. Explain why the natural recursion maintains the correctness of the accumulator statement:

```
(sum/a (rest alon) (+ (first alon) a))
```

Study the above examples before you formulate a general argument. ■

Exercise 407. Complete the above manual evaluation of

```
(sum/a '(10 4 6))
```

Doing so shows that the `sum` and `sum.v2` add up the given numbers in reverse order. While `sum` adds up the numbers from right to left, the accumulator-style version adds them up from left to right. ■

Note For exact numbers, this difference has no effect on the final result. For inexact numbers, the difference is significant. Consider the following definition:

```

(define (g-series n)
  (cond
    [(zero? n) '()]
    [else (cons (expt -0.99 n) (g-series (sub1 n)))]))

```

Applying `g-series` to a natural number produces the beginning of a so-called geometric series. The numbers in this series rapidly oscillate in the interval $(-1, +1)$:

```

> (g-series 5)

(list
  #i-0.9509900498999999
  #i0.96059601
  #i-0.970299
  #i0.9801
  #i-0.99)

```

Each summation function computes a different sum for the same series:

```

> (sum (g-series 1000.0))
#i-0.49746596003269394
> (sum.v2 (g-series 1000.0))
#i-0.49746596003269533

```

While the difference may appear to be small, imagine a context where the actual calculation is something like this:

```

> (- (* 1e+16 (sum (g-series 1000.0)))
     (* 1e+16 (sum.v2 (g-series 1000.0))))
#i14.0

```

And now the difference matters. **End**

For the second example, we turn to the well-known factorial function:

```
; N -> N
; compute (* n (- n 1) (- n 2) ... 1)

(check-expect (! 3) 6)

(define (! n)
  (cond
    [(zero? n) 1]
    [else (* n (! (sub1 n)))]))
```

While `relative-2-absolute` and `invert` processed lists, the factorial function works on natural numbers. Its template is that for `N` processing functions.

We proceed as before by with a template for an accumulator-style version:

```
; N -> N
; compute (* n0 (- n0 1) (- n0 2) ... 1)

(check-expect (!.v2 3) 6)

(define (!.v2 n0)
  (local (; N ??? -> N
        ; compute (* n (- n 1) (- n 2) ... 1)
        ; accumulator ...
        (define (!/a n a)
          (cond
            [(zero? n) ...]
            [else (... (!/a (sub1 n) ... a ...))]))
    (!/a n0 ...)))
```

followed by a sketch of a hand evaluation:

<pre>(! 3) = = (* 3 (! 2)) = (* 3 (* 2 (! 1))) = (* 3 (* 2 (* 1 (! 0)))) = (* 3 (* 2 (* 1))) ... = 6</pre>	<pre>(!.v2 3) = = (!/a 3 a0) = (!/a 2 ... 3 ... a0) = (!/a 1 ... 2 ... 3 ... a0) = (!/a 0 ... 1 ... 2 ... 3 ... a0) ... = 6</pre>
--	---

The left column shows how the original version works, the right one sketches how the accumulator-style function proceeds. Both traverse the natural number until they reach 0. While the original version schedules only multiplications, the accumulator keeps track of each number as the structural processing descends through the given natural number.

Given the goal of multiplying these numbers, `!/a` can use the accumulator to multiply the numbers immediately:

a is the product of the natural numbers in the interval $[n_0, n]$.

In particular, when `n0` is 3 and `n` is 1, a is 6.

Exercise 408. What should the value of a be when `n0` is 3 and `n` is 1? How about when `n0` is 10 and `n` is 8? ■

The factorial function is useful in certain areas of mathematics and in some programming libraries. No ordinary programmer has a need to know the function for its own sake.

Using this invariant we can easily pick the initial value for `a`—it is `1`—and we know that multiplication the current accumulator with `n` is the proper update operation:

```
(define (!.v2 n0)
  (local (; N N -> N
          ; compute (* n (- n 1) (- n 2) ... 1)
          ; accumulator a is the product of the natural
          ; numbers in the interval [n0,n).
          (define (!/a n a)
            (cond
              [(zero? n) a]
              [else (!/a (sub1 n) (* n a))])))
    (!/a n0 1)))
```

It also follows from the accumulator statement that when `n` is `0`, the accumulator is the product of `n` through `1`, meaning it is the desired result. So, like `sum`, `!/a` returns `a` in this case and uses the result of the recursion in the second case.

Exercise 409. Like `sum`, `!` performs the primitive computation steps—multiplication in this case—in reverse order. Surprisingly, this affects the performance of the function in a negative manner.

Measure how long it takes to evaluate `(! 20)` one thousand times. Recall that `(time an-expression)` function determines how long it takes to run `an-expression`. ■

For the third and last example, we use a function that measures the height of simplified binary trees. The example illustrates that accumulator-style programming applies to all kinds of data, not just those defined with single self-references. Indeed, it is as common for complicated data definitions as it is for lists and natural numbers.

Here are the relevant definitions:

```
(define-struct node (left right))
; A BinaryTree (short Tree) is one of:
; - '()
; - (make-node Tree Tree)

(define example (make-node (make-node '() (make-node '() '())) '()))
```

These trees carry no information; their leaves are `'()`. Still, there are many different trees as [figure 108](#) shows; it also uses suggestive graphics to bring across what these pieces of data look like as trees.


One property one may wish to compute is the height of such a tree:

```
; Tree -> N
; measure the height of abt0

(check-expect (height example) 3)

(define (height abt)
  (cond
    [(empty? abt) 0]
    [else (+ (max (height (node-left abt)) (height (node-right abt))) 1)]))
```

The table in [figure 108](#) indicates how to measure the height of a tree though it leaves the notion somewhat ambiguous: it is either the number of nodes from the root of the tree to the highest leaf or the number of connections on such a path. The `height` function follows the second option.



To transform this function into an accumulator-style function, we follow the standard path. We begin with an appropriate template:

```
; Tree -> N
; measure the height of abt0

(check-expect (height.v2 example) 3)

(define (height.v2 abt0)
  (local (; Tree ??? -> N
          ; measure the height of abt
          ; accumulator ...
          (define (height/a abt a)
            (cond
              [(empty? abt) ...]
              [else
               (... (height/a (node-left abt) ... a ...) ...
                    ... (height/a (node-right abt) ... a ...) ...)])))
    (height/a abt0 ...)))
```

As always, the problem is to determine what knowledge the accumulator represents. One obvious choice is the number of traversed branches:

a is the number of steps it takes to reach abt from $abt0$.

Illustrating this accumulator invariant is best done with a graphical example. Take a second look at [figure 108](#). The bottom-most tree comes with two annotations, each pointing out one subtree:

1. If $abt0$ is the complete tree and abt is the subtree pointed to by the circled 1, the accumulator's value must be 1 because it takes exactly one step to get from the root of abt to the root of $abt0$.
2. In the same spirit, for the subtree labeled 2 the accumulator is 2 because it takes two steps to get this place.

As for the preceding two examples, the invariant basically dictates how to follow the rest of the design recipe for accumulators: the initial value for `a` is `0`; the update operation is `add1`; and the base case uses the accumulated knowledge by returning it. Translating this into code yields the following skeleton definition:

```
(define (height.v2 abt0)
  (local (; Tree N -> N
          ; measure the height of abt
          ; accumulator a is the number of steps
          ; it takes to reach abt from abt0

          (define (height/a abt a)
            (cond
              [(empty? abt) a]
              [else
               (... (height/a (node-left abt) (+ a 1)) ...
                    ... (height/a (node-right abt) (+ a 1)) ...)])))
    (height/a abt0 0)))
```

But, in contrast to the first two examples, `a` is not the final result. In the second `cond` clause, the two recursive calls yield two values. The design recipe for structural functions dictate that we combine those in order to formulate an answer for this case; the dots above indicate that we still need to pick an operation that combines these values.

Following the design recipe also tells us that we need to interpret the two values to find the appropriate function. According to the purpose statement for `height/a`, the first value is the height of the left subtree, and the second one is the height of the right one. Given that we are interested in the height of `abt` itself and that the height is the largest number of steps it takes to reach a leaf, we use the `max` function to pick the proper one; see [figure 109](#) for the complete definition.

```
; Tree -> N
; measure the height of abt0

(check-expect (height.v2 example) 3)

(define (height.v2 abt0)
  (local (; Tree N -> N
          ; measure the height of abt
          ; accumulator a is the number of steps
          ; it takes to reach abt from abt0

          (define (height/a abt a)
            (cond
              [(empty? abt) a]
              [else
               (max (height/a (node-left abt) (+ a 1))
                    (height/a (node-right abt) (+ a 1)))])))
    (height/a abt0 0)))
```

Figure 109: The accumulator-style version of *height*

Note on an Alternative Design In addition to counting the number of steps it takes to reach a node, an accumulator function could hold on to the largest height encountered so far. Here is the accumulator statement for the design idea:

the first accumulator, `a`, represents the number of steps it takes to reach `abt`

from `abt0` and the second accumulator, stands for the tallest branch in the part of `abt0` that is to the left of `abt`.

Clearly, this statement assumes a template with two accumulator parameters:

```
(define (height.v3 abt0)
  (local (; Tree N N -> N
          ; measure the height of abt
          ; accumulator s is the number of steps
          ; it takes to reach abt from abt0
          ; accumulator m is the maximal height of
          ; the part of abt0 that is to the left of abt
          (define (h/a abt s m)
            (cond
              [(empty? abt) ...]
              [else
               (... (h/a (node-left abt) ... s ... m ...) ...
                    ... (h/a (node-right abt) ... s ... m ...) ...)])))
    (h/a abt0 ...)))
```

Exercise 410. Complete the design of `height.v3`.

Hint In terms of the bottom-most tree of [figure 108](#), the place marked 1 has no complete paths to leaves to its left while the place marked 2 has one complete path and it consists of two steps. ■

This second design has a more complex accumulator invariant than the first one. By implication, its implementation requires more care than the first one. At the same time, it comes without any advantages, meaning it is inferior to the first one.

Our point is that different accumulator invariants yield different variants. You can design both variants systematically following the same design recipe. When you have complete function definitions, you can compare and contrast the results, and you can then decide which one to keep based on evidence. **End**

Exercise 411. Design an accumulator-style version of `product`, the function that computes the product of a list of numbers. Stop when you have formulated the accumulator invariant and have someone check it. ■

Exercise 412. Design an accumulator-style version of `how-many`, which is the function that determines the number of items on a list. Stop when you have formulated the accumulator invariant and have someone check it. ■

Exercise 413. Design an accumulator-style version of `add-to-pi`, which adds a natural number to `pi` without using `+`:

```
; N -> Number
; add n to pi without use +

(check-within (add-to-pi 2) (+ 2 pi) 0.001)

(define (add-to-pi n)
  (cond
    [(zero? n) pi]
    [else (add1 (add-to-pi (sub1 n)))])))
```

Stop when you have formulated the accumulator invariant and have someone check it. ■

Exercise 414. Design the function `make-palindrome`, which accepts a non-empty list and

constructs a palindrome by mirroring the list around the last item. When given `(explode "abc")`, it yields `(explode "abcba")`.

Hint Here is a solution designed by function composition:

```
; [cons 1String [List-of 1String]] -> [cons 1String [List-of 1String]]
; create a palindrome by mirroring the given list around the last item

(check-expect (mirror (explode "abc")) (explode "abcba"))

(define (mirror s0)
  (append (all-but-last s0) (list (last s0)) (reverse (all-but-last s0))))
```

See [Generalizing Functions](#) for `last`; design `all-but-last` in an analogous manner. This solution traverses `s0` four times:

1. via `all-but-last`,
2. via `last`,
3. via `all-but-last` again, and
4. via `reverse`, which is ISL+'s version of inverse.

Even with `local` definition for the result of `all-but-last`, the function needs three traversals. While these traversals aren't "stacked" and therefore don't have a disastrous impact on the function's performance, an accumulator version can compute the same result with a single traversal. ■

Exercise 415. [Exercise 384](#) implicitly asks for the design of a function that rotates a [Matrix](#) until the first coefficient of the first row differs from 0. In the context of [Exercise 384](#), the solution calls for a generative recursive function that creates a new matrix by shifting the first row to the end when it encounters a 0 in the first position. Here is the solution:

```
; Matrix -> Matrix
; find the first row that doesn't start with 0 and use it as the first one
; generative move the first row to last place
; termination this function does not terminate if all rows in M start with 0

(check-expect (rotate-until.v2 '((0 4 5) (1 2 3))) '((1 2 3) (0 4 5)))

(define (rotate M)
  (cond
    [(not (= (first (first M)) 0)) M]
    [else (rotate (append (rest M) (list (first M))))]))
```

Stop! Modify this function so that it signals an error when all rows start with 0.

If you measure this function on large instances of [Matrix](#), you get a surprising result:

rows in M	1000	2000	3000	4000	5000
<code>rotate</code>	17	66	151	272	436

As the number of rows increases from 1,000 to 5,000, the time spent by `rotate` does not increase by a factor of five but by twenty.

The problem is that `rotate` uses `append`, which makes a brand-new list like `(rest M)` only to add `(first M)` at the end. If `M` consists of 1,000 rows and the last row is the only

one with a non-0 coefficient, that's roughly

$$1,000 \cdot 1,000 = 1,000,000$$

lists. How many lists do we get if M consists of 5,000 lines?

Now suppose we conjecture that the accumulator-style version is faster than the generative one. Here is the accumulator template assuming `rotate` is a structurally recursive function:

```
; Matrix -> Matrix
; find the first row that doesn't start with 0 and use it as the first one
(define (rotate.v2 M0)
  (local (; Matrix ... -> Matrix
          ; accumulator ...
          (define (rotate/a M seen)
            (cond
              [(empty? M) ...]
              [else (... (rotate/a (rest M) ... seen ...) ...)])))
    (rotate M0 ...)))
```

The goal is to remember the first row when its leading coefficient is 0 without using `append` for every recursion.

Formulate an accumulator statement. Then follow the accumulator design recipe to complete the above function. Measure how fast it runs on a `Matrix` that consists of rows with leading 0s except for the last one. If you completed the design correctly, the function is quite fast. ■

Exercise 416. Design `to10`. It consumes a list of digits and produces the corresponding number. The first item on the list is the **most significant** digit. Hence, when applied to `'(1 0 2)`, it produces `102`.

Domain Knowledge You may recall from grade school that the result is determined by

$$1 \cdot 10^2 + 0 \cdot 10^1 + 2 \cdot 10^0 = ((1 \cdot 10 + 0) \cdot 10) + 2 = 102.$$

■

Exercise 417. Design the function `is-prime`, which consumes a natural number and returns `#true` if it is prime and `#false` otherwise.

Domain Knowledge A number n is prime if it is not divisible by any number between $n - 1$ and 2.

Hint The design recipe for `N [≥ 1]` suggests the following template:

```
; N [≥ 1] -> Boolean
; determine whether n is a prime number
(define (is-prime? n)
  (cond
    [(= n 1) ...]
    [else (... (is-prime? (sub1 n)) ...)]))
```

This template immediately tells you that the function forgets n , its initial argument as it recurs. Since n is definitely needed to determine whether n is divisible by `(- n 1)`, `(- n 2)`, and so on, you know that you need an accumulator-style function. ■

Note on Speed People who encounter accumulator-style programming for the first time often get the impression that they are always faster than their recursive counterparts.

Both parts are plain wrong. While it is impossible to explain this mistake in reasoning in this book, let us take a look at the solution of [exercise 409](#):

!	5.760	5.780	5.800	5.820	5.870	5.806
!.v2	5.970	5.940	5.980	5.970	6.690	6.111

The table represents the timings for the two factorial functions. Specifically, the top row shows the number of seconds for 1,000 evaluations of `(! 20)` where the last cell shows the average. The bottom row shows the result of an analogous experiment with `!.v2`. Bottom line is the performance of the accumulator-style version of factorial is always worse than that of the original factorial function.

37.4 A Graphical Editor, with Mouse

[A Graphical Editor](#) introduces the notion of an one-liner editor and presents a number of exercises on creating a graphical editor. An graphical editor is an interactive program that interprets key events as editing actions on a string. In particular, when a user presses the left or right arrow key, the cursor moves left or right; similarly, pressing the delete key removes a `1String` from the edited text. The editor program uses a data representation that combines two strings in a structure. [A Graphical Editor, Revisited](#) resumes these exercises and shows how the same program can greatly benefit from a different data structure, one that combines two strings.

Neither of these sections deals with mouse actions for navigation, even though all modern applications support this functionality. The basic difficulty with mouse events is to place the cursor at the appropriate spot. Since the program deals with a single line of text, a mouse click at (x,y) clearly aims to place the cursor between the letters that are visible at or around the x position. This section fills the gap.

Recall the relevant definitions from [A Graphical Editor, Revisited](#):

```
(define FONT-SIZE 11)
(define FONT-COLOR "black")

; [List-of 1String] -> Image
; render a string as an image for the editor
(define (editor-text s)
  (text (implode s) FONT-SIZE FONT-COLOR))

(define-struct editor [pre post])
; An Editor is
;   (make-editor [List-of 1String] [List-of 1String])
; interpretation if (make-editor p s) is the state of an
; interactive editor, (reverse p) corresponds to the text to the
; left of the cursor and s to the one on its right
```

Exercise 418. Use the structural design recipe to develop `split-structural`. The function consumes a list of `1Strings` `e` and a natural number `x`; the former represents the complete string in some `Editor` and the latter the x coordinate of the mouse click. It produces

```
(make-editor p s)
```

such that (1) `p` and `s` make up `e` and (2) `x` is larger than the image of `p` and smaller than the image of `p` extended with the first `1String` on `s` (if any).

Here is the first condition expressed with an ISL+ expression:

```
(string=? (string-append p s) e)
```

The second one is

```
(<= (image-width (editor-text p))
     x
     (image-width (editor-text (append p (first s))))))
```

assuming `(cons? s)`.

Hints (1) The x coordinate measures the distance from the left. Hence the function must check whether smaller and smaller prefixes of e fit into the given width. The first one that doesn't corresponds to the `pre` field of the desired [Editor](#), the remainder of e to the `post` field.

(2) Designing this function calls for thoroughly developing examples and tests. See [Intervals](#), [Enumerations](#), [Itemizations](#). ■

Exercise 419. Design the function `split`. Use the accumulator design recipe to improve on the result of [exercise 418](#). After all, the hints already point out that when the function discovers the correct split point, it needs both parts of the list and one part is obviously lost due to recursion. ■

Once you have solved this exercise, equip the `main` function of [A Graphical Editor](#), [Revisited](#) with a clause for mouse clicks. As you experiment with moving the cursor via mouse clicks, you will notice that it does not exactly behave like applications that you use on your other devices—even though `split` passes all its tests.

Graphical programs, like editors, call for experimentation to come up with best “look and feel” experiences. In this case, your editor is too simplistic with its placement of the cursor. After the applications on your computer determine the split point, they also determine which letter division is closer to the x coordinate and place the cursor there.

38 More Uses of Accumulation

This chapter presents three more uses of accumulators. The first section concerns the use of accumulators in conjunction with tree-processing functions. It uses the compilation of ISL+ as an illustrative example. The second section explains why we occasionally want accumulators inside of data representations and how to go about it. The final section resumes the discussion of rendering fractals.

38.1 Accumulators and Trees

When you ask DrRacket to run an ISL+ program, it translates the program to commands for your specific computer. This process is called *compilation* and the part of DrRacket that performs the task is called a *compiler*. Before the compiler translates the ISL+ program, it checks that every variable is declared via a `define`, `define-struct`, or a `lambda`.

Stop! Enter x , `(lambda (y) x)`, and `(x 5)` as completes ISL+ programs into DrRacket and ask it to run each. What do you expect to see?

Let's phrase this idea as a sample problem:

Sample Problem: You have been hired to re-create a part of the ISL+ compiler. Specifically, your task deals with the following language fragment, specified in the so-called grammar notation that many programming language manuals use:

```
expression = variable
            | (λ (variable) expression)
            | (expression expression)
```

Remember from [BSL: Grammar](#) that you can read the grammar aloud replacing = with “is one of” and | with “or.”

Recall that λ expressions are functions without names. They bind their parameter in their body. Conversely, a variable occurrence is declared by a surrounding λ that specifies the same name as a parameter. You may wish to revisit [Intermezzo: Scope](#) because it deals with the same issue from the perspective of a programmer. Look for the terms “binding occurrence,” “bound occurrence,” and “free.”

Develop a data representation for the above language fragment; use symbols to represent variables. Then design a function that replaces all undeclared variables with `'*undeclared`.

This problem is representative of many steps in the translation process and, at the same time, is a great case study for accumulator-style functions.

Before we dive into the problem, let's look at some examples in this mini-language, recalling what we know about [lambda](#):

- $(\lambda (x) x)$ is the function that returns whatever it is given, also known as the identity function;
- $(\lambda (x) y)$ looks like a function that returns y whenever it is given an argument, except that y isn't declared;
- $(\lambda (y) (\lambda (x) y))$ is a function that, when given some value v , produces a function that always returns v ;
- $((\lambda (x) x) (\lambda (x) x))$ applies the identity function to itself;
- $((\lambda (x) (x x)) (\lambda (x) (x x)))$ is a short infinite loop; and
- $((((\lambda (y) (\lambda (x) y)) (\lambda (z) z)) (\lambda (w) w)))$ is complex expression that is best run in ISL+ to find out whether it even terminates.

Indeed, you can run all of the above ISL+ expression in DrRacket to confirm what is written about them.

Exercise 420. Explain the scope of each binding occurrence in the above examples. Draw arrows from all bound occurrences to the binding occurrences. ■

Developing a data representation for the language is easy, especially because its description uses a grammar notation. Here is one possibility:

```
; A Lam is one of:
```

We use the Greek letter λ instead of [lambda](#) to signal that this exercise deals with ISL+ as an object of study not just a programming language.


```

; - a Symbol
; - (list 'λ (list Symbol) Lam)
; - (list Lam Lam)

```

Because of `quote`, this data representation makes it easy to create data representations for expressions in our subset of ISL+:

```

(define ex1 '(λ (x) x))
(define ex2 '(λ (x) y))
(define ex3 '(λ (y) (λ (x) y)))
(define ex4 '((λ (x) (x x)) (λ (x) (x x))))

```

These four data examples are representations of some of the above expressions. Stop! Create data representations for the remaining examples.

Exercise 421. Define the functions `is-var?`, `is-λ?`, and `is-app?`, predicates that distinguish (representations of) variables from λ expressions and applications.

Also define

- `λ-para`, which extracts the parameter from a λ expression;
- `λ-body`, which extracts the body from a λ expression;
- `app-fun`, which extracts the function from an application; and
- `app-arg`, which extracts the argument from an application.

With these, you basically can act as if you had used a structure-oriented data definition.

Design `declareds`, which produces the list of all symbols used as λ parameters in a λ term. Don't worry about duplicate symbols. ■

Exercise 422. Develop a data representation for the same subset of ISL+ that uses structures instead of lists. Also provide data representations for `ex1`, `ex2`, and `ex3` following your data definition. ■

We follow the structural design recipe, and here is the product of the steps two and three:

```

; Lam -> Lam
; replace all symbols s in le with '*undeclared if they do
; not occur within the body of a λ whose parameter is s

(check-expect (undeclareds ex1) ex1)
(check-expect (undeclareds ex2) '(λ (x) *undeclared))
(check-expect (undeclareds ex3) ex3)
(check-expect (undeclareds ex4) ex4)

(define (undeclareds le0)
  le0)

```

Note how we expect `undeclareds` to process `ex4` even though the expression loops forever when run; compilers don't run programs, they read them and create others.

A close look at the purpose statement directly suggests that the function needs an accumulator. This becomes even clearer when we inspect the template for `undeclareds`:

```

(define (undeclareds le)
  (cond

```

```

[[is-var? le) ...]
[[is-λ? le) (... (undeclareds (λ-body le)) ...)]
[[is-app? le)
 (... (undeclareds (app-fun le)) (undeclareds (app-arg le)) ...)]))

```

When undeclareds recurs on the body of (the representation of) a λ expression, it forgets $(\lambda\text{-para } le)$, the declared variable.

So, let's start with an accumulator-style template:

```

(define (undeclareds le0)
  (local (; Lam ??? -> Lam
          ; accumulator a represents ...
          (define (undeclareds/a le a)
            (cond
              [(is-var? le) ...]
              [(is-λ? le) (... (undeclareds (λ-body le) ... a ...) ...)]
              [(is-app? le)
               (... (undeclareds (app-fun le) ... a ...)
                    (undeclareds/a (app-arg le) ... a ...) ...)])))
    (undeclareds/a le0 ...)))

```

In this context, we can now formulate an accumulator invariant:

a represents the list of λ parameters encountered on the path from le_0 to le .

For example, if le_0 is

```
'(((λ (y) (λ (x) y)) (λ (z) z)) (λ (w) w))
```

and le is the highlighted subtree, then a contains y . The left side of figure 110 presents a graphical illustration of the same example. It shows a Lam expression as an upside-down tree, that is, the root is at the top. A $@$ node represents an application with two descendants; the other nodes are self-explanatory. In this tree diagram, the bold path leads from le_0 to le through a single variable declaration.

Similarly, if we pick a different subtree of the same piece of data,

```
'(((λ (y) (λ (x) y)) (λ (z) z)) (λ (w) w))
```

we get an accumulator that contains both y and x . The right side of figure 110 makes this point again. Here the bold path leads through two λ nodes to the boxed subtree, and the accumulator is the list of declared variables along the bold path.




Figure 110: Lam terms as trees

Now that we have settled on the data representation of the accumulator and its

invariant, we can resolve the remaining design questions:

- the initial accumulator value of '() ;
- we can use `cons` to add $(\lambda\text{-para } le)$ to `a`; and
- once `undeclareds/a` reaches a variable, it can use the accumulator to check whether the variable is in the scope of a declaration.

Figure 111 shows how to translate these ideas into a complete function definition. Note the name `declareds` for the accumulator; it brings across the key idea behind the accumulator invariant, helping the programmer understand the definition. The base case uses `member?` from ISL+ to determine whether the variable `le` is in `declareds` and, if not, replaces it with `'*undeclared`. The second `cond` clause uses a `local` to introduce the extended accumulator `newd`. Because `para` is also used to rebuild the expression, it has its own local definition. Finally, the last clause concerns function applications, which do not declare variables and do not use any directly. As a result, it is by far the simplest of the three clauses.

```
; Lam -> Lam
; replace all symbols s in le with '*undeclared if they do
; not occur within the body of a λ whose parameter is s

(check-expect (undeclareds ex1) ex1)
(check-expect (undeclareds ex2) '(\ (x) *undeclared))
(check-expect (undeclareds ex3) ex3)
(check-expect (undeclareds ex4) ex4)

(define (undeclareds le0)
  (local (; Lam [List-of Symbol] -> [List-of Symbol]
        ; accumulator declareds is a list of all λ
        ; parameters on the path from le0 to le
        (define (undeclareds/a le declareds)
          (cond
            [(is-var? le)
             (if (member? le declareds) le '*undeclared)]
            [(is-λ? le)
             (local ((define para (λ(para le))
                               (define newd (cons para declareds))
                               (define body (undeclareds/a (λ-body le) newd)))
               (list 'λ (list para body)))]
            [(is-app? le)
             (list (undeclareds/a (app-fun le) declareds)
                   (undeclareds/a (app-arg le) declareds))]))
        (undeclareds/a le0 '())))
```

Figure 111: Finding undeclared variables

Exercise 423. Make up an ISL+ expression in which `x` occurs both free and bound. Formulate it as an element of `Lam`. Does `undeclareds` work properly on your expression? ■

Exercise 424. Considering the following expression:

```
(λ (*undeclared) ((λ (x) (x *undeclared)) y))
```

Yes, it uses `*undeclared` as a variable. Represent it in `Lam` and check what `undeclareds`

produces for this expression.

Modify `undeclareds` so that it replaces each free occurrence of `'x` with

```
(list '*undeclared 'x)
```

and each bound one `y` with

```
(list '*declared 'y)
```

Doing so unambiguously identifies problem spots, which a program-development environment such as DrRacket can use to high-light errors.

Note The trick to replace a variable occurrence with the representation of an application feels awkward. If you dislike it, consider synthesizing the symbols `'*undeclared:x` and `'declared:y` instead. ■

Exercise 425. Re-design the `undeclareds` function for the structure-based data representation from [exercise 422](#). ■

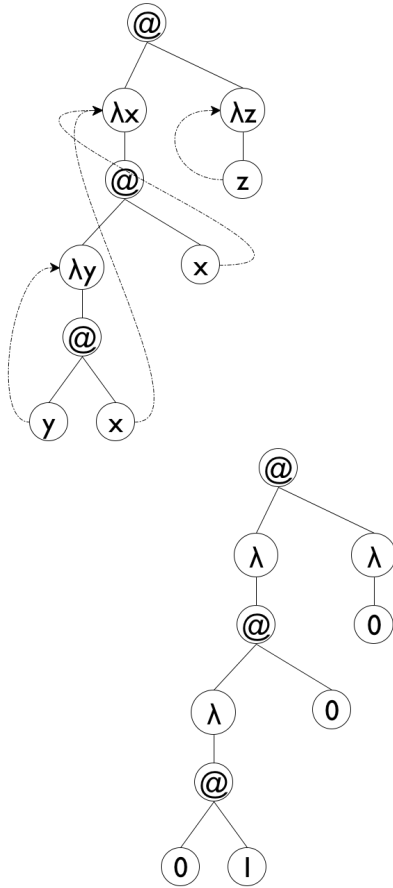


Figure 112: Static distances

Exercise 426. Design the function `static-distance`. It replaces all occurrences of variables with a natural number that represents how far away the declaring `λ` is.

[Figure 112](#) illustrates the idea. The tree on the left represents the [Lam](#) term

```
'((λ (y) (λ (x) (y x))) (λ (z) z))
```

in graphical form. It includes dotted arrows that point from variable occurrences to the corresponding variable declarations. On the right, the figure shows a tree of the same shape, though without the arrows. The `'λ` nodes come without names, and variable occurrences have been replaced by natural numbers that specify which `'λ` declares the variable. Each natural number n says that the binding occurrence is n steps upwards—toward the root of the `Lam` tree. A value of `0` denotes the first `'λ` on the path to the root, `1` the second one, and so on.

Hint The `undeclareds` accumulator of `undeclareds/a` is a list of all parameters on path from `le` to `le0` **in reverse order**—the last one seen is at the first on the list. ■

~~38.2 Data Representations with Accumulators~~

When you play board games or solve puzzles, you tend to think about your possible moves at every stage. As you get better, you may even imagine the possibilities after this first step. The result is a so-called *game tree*, which is a (part of the) tree of all possible moves that the rules allow.

Programs that generate and search such game trees are ubiquitous in the world of computing:

Sample Problem: Your manager tells you the following story.


“Once upon a time, three cannibals were guiding three missionaries through a jungle. They were on their way to the nearest mission station. After some time, they arrived at a wide river, filled with deadly snakes and fish. There was no way to cross the river without a boat. Fortunately, they found a row boat with two oars after a short search. Unfortunately, the boat was too small to carry all of them. It could barely carry two people at a time. Worse, because of the river’s width someone had to row the boat back.

“Since the missionaries could not trust the cannibals, they had to figure out a plan to get all six of them safely across the river. The problem was that these cannibals would kill and eat missionaries as soon as there were more cannibals than missionaries at some place. Thus our missionaries had to devise a plan that guaranteed that there were never any missionaries in the minority at either side of the river. The cannibals, however, could be trusted to cooperate otherwise. Specifically, they wouldn’t abandon any potential food, just as the missionaries wouldn’t abandon any potential converts.”


While your manager doesn’t assign any specific design task, he wants to explore whether the company can design (and sell) programs that solve such puzzles.

The point of this sample problem isn’t just another form of tree, but an example of where accumulators are best merged with problem data rather than managed separately in a function.

In principle, it is quite straightforward to solve such puzzles by hand. Here is the rough idea. Pick a graphical representation of the problem states. Ours consists of a three-part box: the left one represents the missionaries and the cannibals on the left; the middle combines the river and the boat; and the third part is the right-hand side of the river. Take a look at the following representation of the initial state:



Black circles denote missionaries, white circles cannibals. All of them are on the left-hand river bank. The boat is also on the left side. Nobody is on the right. Here are two more states:



The first one is the final state, where all people and the boat are on the right bank of the river. The second one depicts some intermediate state where two people are on the left with the boat and four people are on the right.




Figure 113: Creating a game tree

Now that you have a way to write down the state of the puzzle, you can think about the possibilities at each stage. Doing so yields a tree of possible moves. Figure 113 sketches the first two and a half layers in such a tree. The left-most state is the initial one.

Because the boat can transport at most two people and must be rowed by at least one, you have five possibilities to explore: one cannibal rows across, two, one missionary and one cannibal go, one missionary, or two missionaries. These possibilities are represented with five arrows going from the initial state to five intermediate states.

For each of these five intermediate states, you can play the same game again. In figure 113 you see how the game continues for the middle (third) one of the new states. Because there are only two people on the right river bank, you see three possibilities: a cannibal goes back, a missionary, or both. Hence three arrows connect the middle state to the three states on the right side of the tree.

If you keep drawing the tree of possibilities—especially if you are doing it systematically—you sooner or later discover the final state. That is, you create the final state with all people on the right river bank by moving a boat load of people from the left river bank to the right one, and nobody is left back.

A close look at figure 113 reveals two problems with this naive approach to generating the tree of possibilities. The first one is the dashed arrow that connects the middle state on the right to the initial state. It indicates that rowing back the two people from the right to the left gets the puzzle back to its initial state, meaning you're starting over,

which is obviously undesirable. The second problem concerns those states with a star in the top-right corner. In both cases, there are more white-circle cannibals than black-circle missionaries on the left river bank, meaning the cannibals would eat the missionaries. Again, the goal is to avoid such states, making these moves undesirable.

One way to turn this puzzle into a program is to design a function that determines whether some final state—here **the** final state—is reachable from some given state. Here is an appropriate function definition:

```
; PuzzleState -> PuzzleState
; determine whether the final state is reachable from the given state
; generative create search tree of possible boat rides
; termination ???

(check-expect (solve initial-puzzle) final-puzzle)

(define (solve state0)
  (local (; [List-of PuzzleState] -> PuzzleState
        ; generative generate the successor states for all intermediate ones
        (define (solve* los)
          (cond
            [(ormap final? los) (first (filter final? los))]
            [else (solve* (create-next-states los))]))
        (solve* (list state0))))
```

The auxiliary function uses generative recursion, generating all new possibilities given a list of possibilities. If one of the given possibilities is a final state, the function returns it.

Clearly, `solve` is quite generic. As long as you define a collection of *PuzzleStates*, a function for recognizing final states, and a function for creating all “successor” states, `solve` can work on your puzzle.

Exercise 427. The `solve*` function generates all states reachable with n boat trips before it looks at states that require $n + 1$ boat trips, even if some of those boat trips return to previously encountered states. Because of this systematic way of traversing the tree, `solve*` cannot go into an infinite loop. Why?

Terminology This way of searching a tree or a graph is dubbed *breadth-first search*. ■

Exercise 428. Develop a data representation for the states of the missionary-and-cannibal puzzle. Like the graphical representation, a state must obviously record the number of missionaries and cannibals on each side of the river plus the location of the boat. After all, these are the properties of the world that change with one boat trip.

The description of *PuzzleStates* calls for a structure type definition. Represent the above initial, intermediate, and final states in your chosen data representation.

Design the function `final?`, which detects whether in a given state all people are on the right river bank.

Design the function `render-mc`, which maps a state of the missionary-and-cannibal puzzle to an image. ■

The problem is that returning the final state says nothing about how the player can get from the initial state to the final one. In other words, `create-next-states` forgets how it gets to the returned states from the given ones. And this situation clearly calls for an accumulator, but at the same time, the accumulated knowledge is best associated with

the individual `PuzzleStates` not `solve*` or any other function.

Exercise 429. Modify the data representation from [exercise 428](#) so that each state records the sequence of states traversed to get there. Use a list of states.

Articulate and write down an accumulator statement with the data definition that explains the additional field.

Do you have to modify `final?` or `render-mc` to work on this revised data representation?

Exercise 430. Design the `create-next-states` function. It consumes lists of missionary-and-cannibal states and generates the list of all those states that a boat ride can reach.

Ignore the accumulator in the first draft of `create-next-states`, but make sure that the function does not generate states where the cannibals can eat the missionaries.

For the second design, update the accumulator field in the state structures and use it to rule out states that have been encountered on the way to the current state. ■

Exercise 431. Exploit the accumulator-oriented data representation to modify `solve`. The revised function produces the list of states that lead from the initial puzzle state to the final one.

Also consider creating a movie from this list, using `render-mc` to generate the images. Use `run-movie` to display the movie. ■

~~38.3 Accumulators as Results~~

Take a second look at [figure 85](#). It displays a Sierpinski triangle and a suggestion how to create such a fractal. Specifically, the images on the right explain one version of the generative idea behind the process of generating a Sierpinski triangle:

The given problem is a triangle. When the triangle is too small to be subdivided any further, the algorithm does nothing; otherwise, it finds the midpoints of its three sides and deals with the three outer triangles recursively.

In contrast, [Fractals, a First Taste](#) shows how to compose Sierpinski triangles algebraically, a process that does not correspond to this description.

Most programmers expect the word “draw” to mean the action of adding a triangle to some canvas. The `scene+line` function from the “image” library make this idea concrete. The function consumes an image `s` and the coordinates of two points and adds a line through these two points to `s`. It is easy to generalize from `scene+line` to `add-triangle` and from there to `add-sierpinski`:

Sample Problem: Design the `add-sierpinski` function. It consumes an image and three `Posns` describing an triangle. It adds to this given image a Sierpinski triangle whose outer perimeter is the given triangle.

Note how this sample problem implicitly refers to the above process description of how to draw a Sierpinski triangle. In other words, we are confronted with a classical generative-recursive problem and we can start with the classic template of generative recursion and the four central design questions:

- The given problem is trivial if the triangle is too small to be subdivided.

- In the trivial case, the function returns the given image.
- Otherwise the process adds the triangle and determines the midpoints of the given triangle's sides. Each “outer” triangle is then processed recursively.
- Each of these recursive steps produces an image. The remaining question is how to combine these images.

It is straightforward to translate these answers into a skeletal ISL+ definition:

```
; Image Posn Posn Posn -> Image
; generative adds the triangle (a, b, c) to s, subdivides the triangle
; into three by taking the midpoints of its sides, and deals
; with the outer triangles recursively until it is too small
(define (add-sierpinski scene0 a b c)
  (cond
    [(too-small? a b c) scene0]
    [else
     (local ((define scene1 (add-triangle scene0 a b c))
              (define mid-a-b (mid-point a b))
              (define mid-b-c (mid-point b c))
              (define mid-c-a (mid-point c a))
              (define scene2 (add-sierpinski scene0 a mid-a-b mid-c-a))
              (define scene3 (add-sierpinski scene0 b mid-b-c mid-a-b))
              (define scene4 (add-sierpinski scene0 c mid-c-a mid-b-c)))
      (... scene0 ... scene1 ... scene2 ... scene3 ...))]))
```

Since each midpoint is used twice, the skeleton use a `local` expression to formulate the generative step in ISL+. The `local` expression introduces the three new midpoints plus three recursive applications of `sierpinski`. The dots in its body suggests a combination of the scenes.

Exercise 432. In the meantime, we can tackle the wish list that implicitly comes with the above skeleton:

```
; Image Posn Posn Posn -> Image
; add the black triangle a, b, c to scene
(define (add-triangle scene a b c)
  scene)

; Posn Posn Posn -> Boolean
; is the triangle a, b, c too small tp be divided further
(define (too-small? a b c)
  #false)

; Posn Posn -> Posn
; determine the midpoint between a and b
(define (mid-point a b)
  a)
```

Design the three functions.

Domain Knowledge (1) For the `too-small?` function it suffices to measure the distance between two points and to check whether it is below some chosen threshold, say, `10`. The distance between (x_0, y_0) and (x_1, y_1) is

$$\sqrt{(x_0 - x_1)^2 + (y_0 - y_1)^2}$$

that is, the distance of $(x_0 - y_0, x_1 - y_1)$ to the origin.




The midpoint between points (x_0, y_0) and (x_1, y_1) is

$$\left(\frac{x_0 + x_1}{2}, \frac{y_0 + y_1}{2} \right)$$

that is, its coordinates are the midpoints between each pair of coordinates, respectively. ■

Now that we have all the auxiliary functions, it is time to return to the problem of combining the three images that are created by the recursive calls. One obvious guess is to use the `overlay` or `underlay` function, but a test at in the interaction area of DrRacket shows that the functions hide the underlying triangles.



Specifically imagine that the three recursive calls produce the following empty scenes enriched with a single triangle in appropriate locations:

```
> scene1

> scene2

> scene3

```

A combination should look like this figure:



But, combining these shapes with `overlay` or `underlay` does not yield this desired shape:

```
> (overlay scene1 scene2 scene3)

> (underlay scene1 scene2 scene3)

```

Indeed, the image library of ISL+ does not support a function that combines these scenes in an appropriate manner.

Let's take a second look at these interactions. If `scene1` is the result of adding the upper triangle to the given scene and `scene2` is the result of adding a triangle on the lower left, perhaps the second recursive call should add triangles to the result of the first call. Doing so would yield



and handing over this scene to the third recursive call produces exactly what is wanted:



```
; Image Posn Posn Posn -> Image
; generative adds the triangle (a, b, c) to s, subdivides the triangle
; into three by taking the midpoints of its sides, and deals
; with the outer triangles recursively until it is too small
; accumulator the function accumulates the triangles in the given scene
(define (add-sierpinski scene0 a b c)
```

```

(cond
  [(too-small? a b c) scene0]
  [else
    (local ((define scene1 (add-triangle scene0 a b c))
              (define mid-a-b (mid-point a b))
              (define mid-b-c (mid-point b c))
              (define mid-c-a (mid-point c a))
              (define scene2 (add-sierpinski scene1 a mid-a-b mid-c-a))
              (define scene3 (add-sierpinski scene2 b mid-b-c mid-a-b)))
      (add-sierpinski scene3 c mid-c-a mid-b-c))]))

```

Figure 114: Adding Sierpinski triangles to an image

Figure 114 shows the reformulation based on this insight. The three highlights pinpoint the key design idea. All concern the case when the triangle is sufficiently large and it is added to the given scene. Once its sides are sub-divided, the first outer triangle is recursively processed using `scene1`, the result of adding the given triangle. Similarly, the result of this first recursion, dubbed `scene2`, is used for the second recursion, which is about processing the second triangle. Finally, `scene3` flows into the third recursive call. In sum, the novelty is that the accumulator is simultaneously an argument, a medium for collecting knowledge, and the result of the function.

To explore `add-sierpinski` it is best to start from an equilateral triangle and an image that leaves a sufficiently large border. Here are definitions that meet these two criteria:

```

(define MT (empty-scene 400 400))

(define A (make-posn 200 50))
(define B (make-posn 27 350))
(define C (make-posn 373 350))

```

Check what kind of Sierpinski fractal `(add-sierpinski MT A B C)` delivers. Experiment with the definitions from [exercise 432](#) to create sparser and denser Sierpinski triangles than the first one.

Exercise 433. To compute the endpoints of an equilateral Sierpinski triangle, draw a circle and pick three points on the circle that are 120 degrees apart, for example, 120, 240, and 360.

Design the function `circle-pt`:

```

(define CENTER (make-posn 200 200))
(define RADIUS 200)

; Number -> Posn
; determines the point on the circle with CENTER and
; RADIUS whose angle is

; examples
; given: 120/360, what are the x and y coordinates of the desired point
; given: 240/360, what are the x and y coordinates of the desired point
; given: 360/360, what are the x and y coordinates of the desired point

(define (circle-pt factor)
  (make-posn 0 0))

```

Domain Knowledge This design problem calls on knowledge from mathematics. One way to view the problem is as a conversion of a complex number from the polar-

coordinate representation to the `Posn` representation. Read up on `make-polar`, `real-part`, and `imag-part` in ISL+. Another way is to use trigonometry, `sin` and `cos`, to determine the coordinates. If you choose this route, recall that these trigonometry functions compute the sine and cosine in terms of radians, not degrees. Also keep in mind that on-screen positions grow downwards not upwards. ■

Exercise 434. Take a look at the following two images:



They demonstrate how to generate a fractal Savannah tree in the same way that [figure 84](#) shows how to draw a Sierpinski triangle. The image on the left shows what a fractal Savannah tree looks like. The right one explains the generative construction step.

Design the function `add-savannah`. The function consumes an image and four numbers: (1) the x coordinate of a line's base point, (2) the y coordinate of a line's base point, (3) the length of the line, and (4) the angle of the line. It adds a fractal Savannah tree to the given image.

Unless the line is too short, the function adds the specified line to the image. It then divides the line into three segments. It recursively uses the two intermediate points as the new starting points for two lines. The lengths and the angles of the two branches change in a fixed manner, but independently of each other. Use constants to define these changes and work with them until you like your tree well enough.


Hint Experiment with shortening the left branches by at least one third and rotating it left by at least 0.15 degrees. For the right branch, shorten it by at least 20% and rotate it by 0.2 degrees in the opposite direction. ■

Exercise 435. Graphics programmers often need to connect two points with a smooth curve where “smooth” is relative to some perspective. Take a look at these two images:



The left one shows a smooth curve, connecting points A and C ; the right one supplies the perspective point, B , and the angle that an observer would have from this point.

One method for drawing such curves is due to Bézier. It is a prime example of generative recursion, and the following sequence explains the *eureka!* behind the algorithm:



Dr.
Géraldine
Morin
suggested
this
exercise.

Consider the image on the left. It reminds you that the three given points determine a triangle and that the connection between A to C is the focal point of the algorithm. The goal is to pull the line from A to C toward B so that it turns into a smooth curve.

Now turn to the image in the middle. It explains the essential idea of the generative step. The algorithm determines the midpoint on the two observer lines, $A-B$ and $B-C$, as well as the midpoint between these two, $A-B-C$.

Finally, the rightmost image shows how these three new points generate two distinct recursive calls: one deals with the new triangle on the left and the other one with the triangle on the right. More precisely, $A-B$ and $B-C$ become the new observer points and the lines from A to $A-B-C$ and from C to $A-B-C$ become the foci of the two recursive calls.

When the triangle is small enough, we have a trivially solvable case. The algorithm just draws the triangle, and it appears as a point on the given image. You may need to experiment with the notion of “small enough” to make the curve look smooth. ■

39 Summary

This sixth and last part of the book is about designing with accumulators, a mechanism for collecting knowledge during a data structure traversal. Adding an accumulator may fix performance flaws, and for algorithms, it may eliminate termination problems. Your take-away from this part are two and a half design lessons:

1. The first step is to recognize the need for introducing an accumulator. Traversals “forget” pieces of the argument when they step from piece to the next. If you discover that such knowledge could simplify the function’s design, consider introducing an accumulator. The first step is to switch to the **accumulator template**.
2. The second, and key, step is to formulate an accumulator statement. Concisely put, the statement must express **what knowledge** the accumulator gathers **using what kind of data**. Don’t proceed until you have answered these questions. In most cases, the accumulator statement describes the relevant difference between the very original argument and the current argument.
3. The third step, a minor one, is to deduce from the accumulator statement (1) what the initial accumulator value is, (2) how to maintain it during traversal steps, and (3) how to exploit its knowledge.

The idea of accumulating knowledge is ubiquitous, and it appears in many different forms and shapes. It is widely used in so-called functional languages like ISL+. Programmers using imperative languages encounter accumulators in a different way, mostly via assignment statements in primitive looping constructs, because the latter cannot return values. Designing such imperative accumulator programs proceeds just like the design of accumulator functions here, but the details are beyond the scope of this first book on systematic program design.